

# \*\*\* *Diabetes Prediction* \*\*\*

## Step1: Importing the libraries.

```
In [164... import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

## Step 2: Loading a Data

```
In [166... df = pd.read_excel("diabetes_data.xlsx")
```

```
In [167... df
```

```
Out[167... 
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	


768 rows × 9 columns



```
In [168... df.head()
```

Out[168...


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
<b>0</b>	6	148	72	35	0	33.6	C
<b>1</b>	1	85	66	29	0	26.6	C
<b>2</b>	8	183	64	0	0	23.3	C
<b>3</b>	1	89	66	23	94	28.1	C
<b>4</b>	0	137	40	35	168	43.1	2



In [169... `df.tail()`

Out[169...

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
<b>763</b>	10	101	76	48	180	32.9	
<b>764</b>	2	122	70	27	0	36.8	
<b>765</b>	5	121	72	23	112	26.2	
<b>766</b>	1	126	60	0	0	30.1	
<b>767</b>	1	93	70	31	0	30.4	



## Step3: Performing a mapping or transformation on a specific column

In [171... `df.shape`

Out[171... (768, 9)

In [172... `df['Outcome'].value_counts()`

Out[172... Outcome  
 0 500  
 1 268  
 Name: count, dtype: int64

## Step 4: Creating a new binary column named "Outcome." by using mapping

In [174... `df['Outcome'] = df['Outcome'].map({  
 1: 'Diabetic',  
 0: 'Non-Diabetic'  
 })`

In [175... df

Out[175...

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns



In [176... 

```
# separating the data and lables
X =df.drop(columns= 'Outcome',axis = 1)
Y = df['Outcome']
```

In [177... X

Out[177...

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 8 columns



In [178...

Y

Out[178...

0

Diabetic

1

Non-Diabetic

2

Diabetic

3

Non-Diabetic

4

Diabetic

...

763

Non-Diabetic

764

Non-Diabetic

765

Non-Diabetic

766

Diabetic

767

Non-Diabetic

Name: Outcome, Length: 768, dtype: object

In [179...

df.head()

Out[179...

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	C
1	1	85	66	29	0	26.6	C
2	8	183	64	0	0	23.3	C
3	1	89	66	23	94	28.1	C
4	0	137	40	35	168	43.1	2

# Step 5: Creating a histogram using 'plt.hist' to visualize the distribution of the "Outcome" column

In [181...

df.describe()

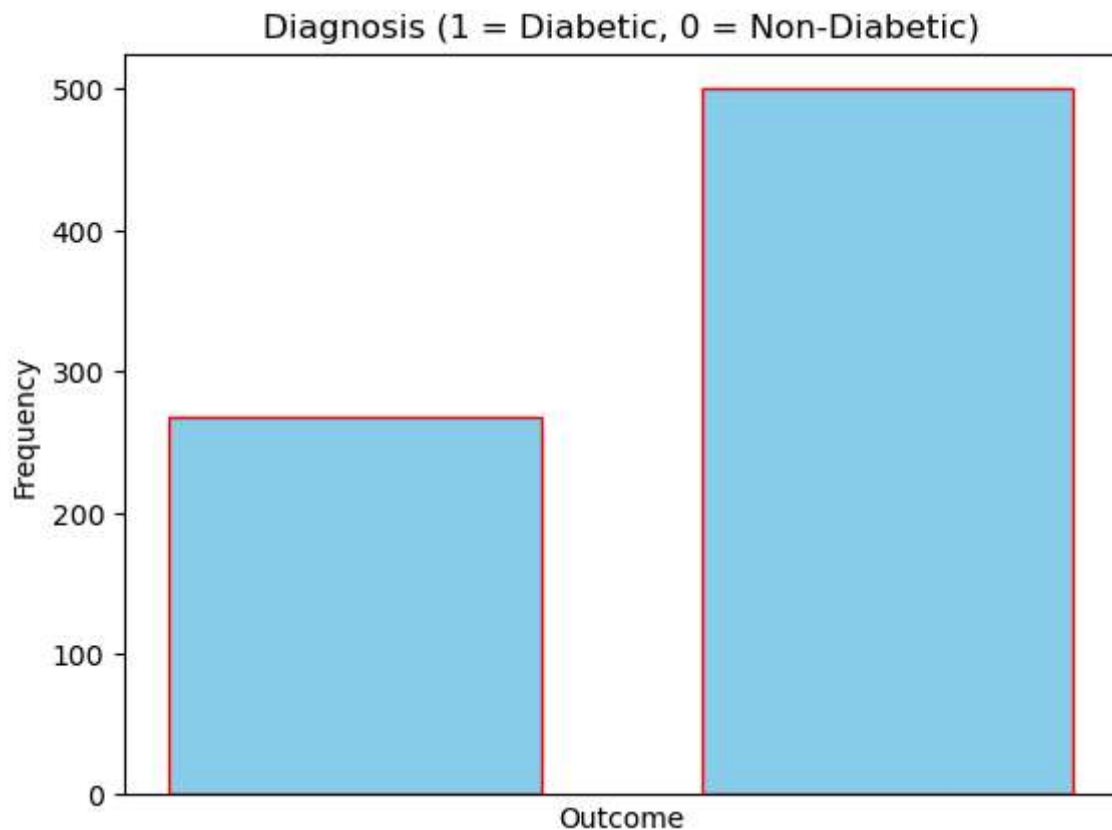
Out[181...

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In [182...

```
# Create histogram for 'Outcome' column
import matplotlib.pyplot as plt

plt.hist(df['Outcome'], bins=2, edgecolor='r', color='skyblue', rwidth=0.7) # Adjust
plt.title('Diagnosis (1 = Diabetic, 0 = Non-Diabetic)')
plt.xlabel('Outcome')
plt.ylabel('Frequency')
plt.show()
```



## step 6: Creating a set of histograms for various features

In [184...

```
import matplotlib.pyplot as plt
import seaborn as sns

# Define features to visualize
X = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age']

# Create subplots
plt.figure(figsize=(10, 12)) # Adjust figure size

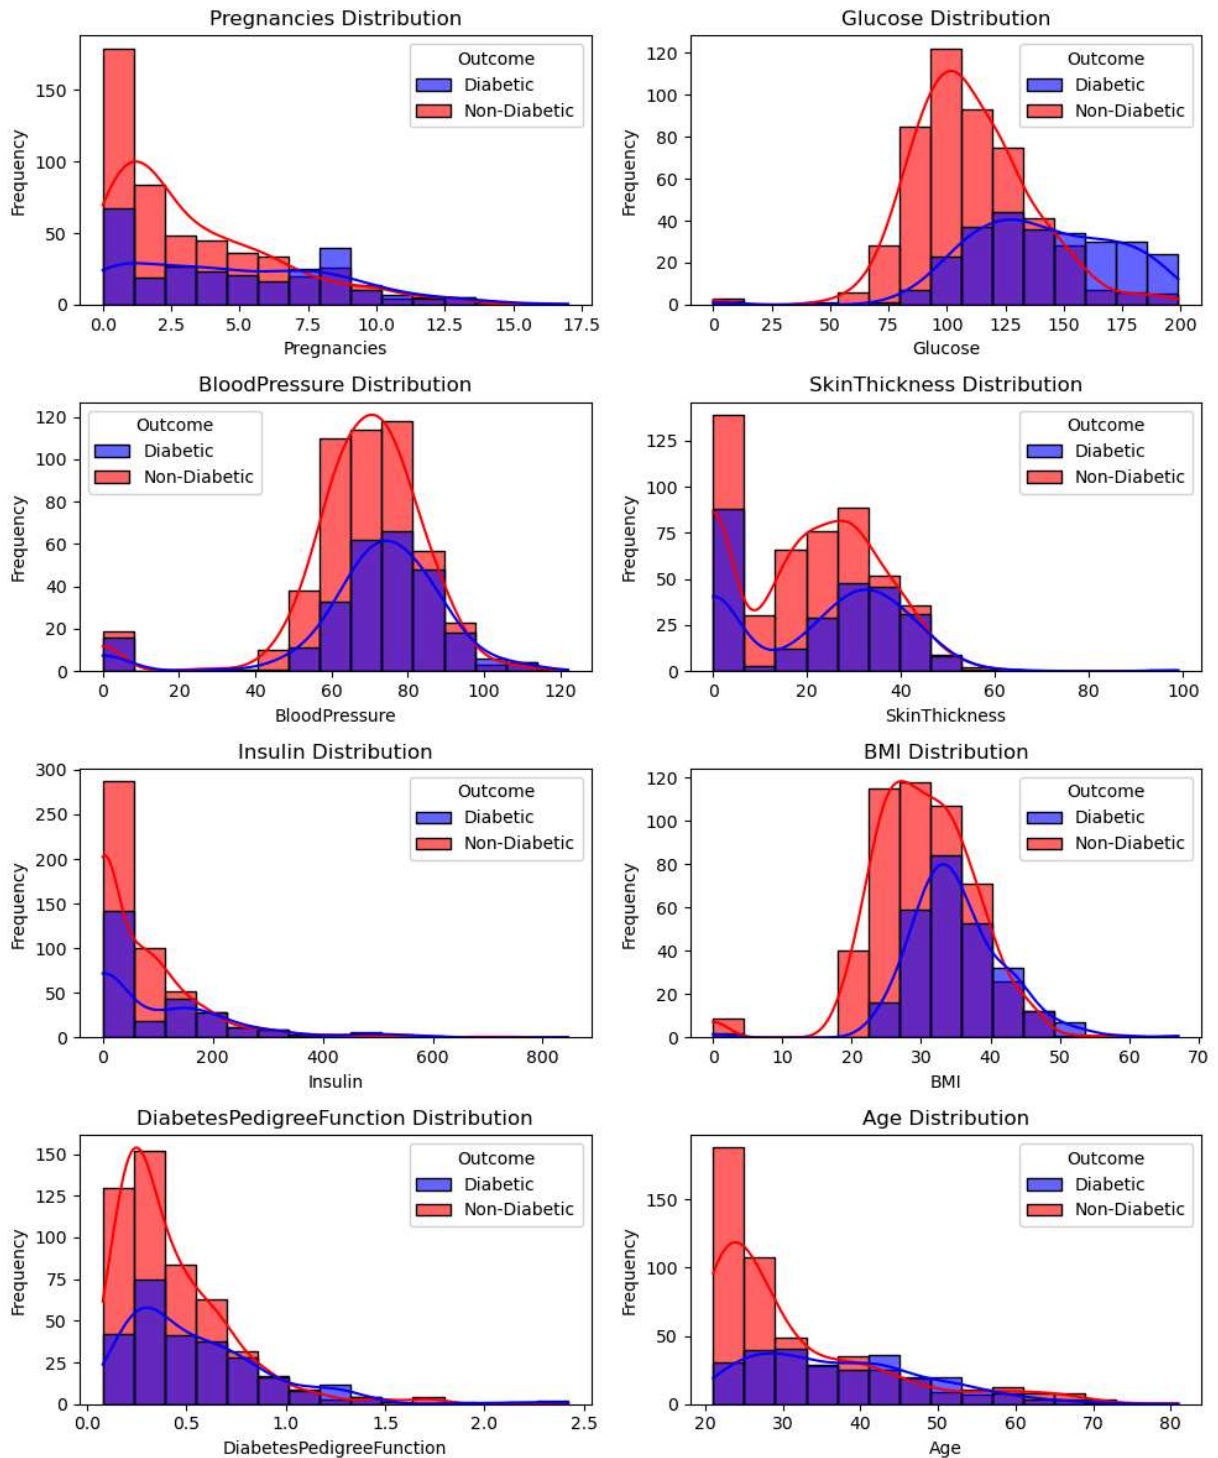
for i, feature in enumerate(X):
    plt.subplot(4, 2, i + 1) # Create subplot grid of 4x2

    # Plot histogram using seaborn for better visualization
    sns.histplot(data=df, x=feature, hue='Outcome', kde=True, bins=15, palette=['b1

    # Add title and labels
    plt.title(f'{feature} Distribution')
    plt.xlabel(feature)
    plt.ylabel('Frequency')

# Adjust layout to prevent overlap
plt.tight_layout()
```

```
# Display the histograms
plt.show()
```



## Step 7: Data Standardising

```
In [208... from sklearn.preprocessing import StandardScaler

# Define the scaler
scaler = StandardScaler()
```

```

# Pass actual numeric data, not column names
standardized_data = scaler.fit_transform(df[X])

# Convert to DataFrame (optional, for better readability)
standardized_df = pd.DataFrame(standardized_data, columns=X)

# Display the first few rows
print(standardized_df.head())

```

```

      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI \
0      0.639947  0.848324      0.149641      0.907270 -0.692891  0.204013
1     -0.844885 -1.123396     -0.160546      0.530902 -0.692891 -0.684422
2      1.233880  1.943724     -0.263941     -1.288212 -0.692891 -1.103255
3     -0.844885 -0.998208     -0.160546      0.154533  0.123302 -0.494043
4     -1.141852  0.504055     -1.504687      0.907270  0.765836  1.409746

      DiabetesPedigreeFunction  Age
0              0.468492  1.425995
1             -0.365061 -0.190672
2              0.604397 -0.105584
3             -0.920763 -1.041549
4              5.484909 -0.020496

```

In [216...

```

X = standardized_data
Y = df['Outcome']
print(X)
print(Y)

```

```

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
0      Diabetic
1    Non-Diabetic
2      Diabetic
3    Non-Diabetic
4      Diabetic
...
763  Non-Diabetic
764  Non-Diabetic
765  Non-Diabetic
766      Diabetic
767  Non-Diabetic
Name: Outcome, Length: 768, dtype: object

```

In [220...

```

df.groupby('Outcome').mean()

```



	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	D
Outcome							
Diabetic	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	
Non-Diabetic	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	

## Step 7: Splitting the data into training and testing

```
In [232...] X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, random_sta
```

```
In [234...] train_df, test_df = train_test_split(df)
# Print the sizes of the training and testing sets
print(f'Training set size: {len(train_df)} rows')
print(f'Testing set size: {len(test_df)} rows')
```

Training set size: 576 rows

Testing set size: 192 rows

## Step 8: Training the model by using Support Vector Machine

```
In [237...] model = svm.SVC(kernel = 'linear')
```

```
In [239...] model.fit(X_train,Y_train)
```

```
Out[239...] SVC
SVC(kernel='linear')
```

```
In [241...] X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction,Y_train)
```

```
In [243...] print(f'Training Accuracy: {training_data_accuracy:.2f}')
```

Training Accuracy: 0.78

```
In [245...] X_test_prediction = model.predict(X_test)
testing_data_accuracy = accuracy_score(X_test_prediction,Y_test )
```

```
In [247...] print(f'Testing Accuracy: {testing_data_accuracy:.2f}')
```

Testing Accuracy: 0.75

## Step 9: Classification model

```
In [250... def classification_model(model, X_train, X_test, Y_train, Y_test, n_folds=5):  
    # Train the model  
    model.fit(X_train, Y_train)  
  
    # Predict and evaluate  
    predictions = model.predict(X_test)  
    accuracy = accuracy_score(Y_test, predictions)  
    print(f'Testing Accuracy: {accuracy:.2f}')  
    # Cross-validation  
    cv_scores = cross_val_score(model, X_train, Y_train, cv=n_folds)  
    print(f'Mean Cross-Validation Score: {cv_scores.mean():.2f}')  
    classification_model(model, X_train, X_test, Y_train, Y_test, n_folds=5)
```

Testing Accuracy: 0.75

Mean Cross-Validation Score: 0.78

## Step 10: Logistic Regression model

```
In [253... logistic_model = LogisticRegression()  
classification_model(logistic_model, X_train, X_test, Y_train, Y_test)  
# Create a Logistic Regression model  
model = LogisticRegression() # Increase max_iter for potential convergence issu  
# Train the model  
model.fit(X_train, Y_train)  
# Make predictions on the test set  
y_pred = model.predict(X_test)  
# Evaluate the model  
accuracy = accuracy_score(Y_test, y_pred)  
print(f"Accuracy: {accuracy}")  
# Print the confusion matrix  
conf_matrix = confusion_matrix(Y_test, y_pred)  
print("Confusion Matrix:")  
print(conf_matrix)
```

Testing Accuracy: 0.74

Mean Cross-Validation Score: 0.77

Accuracy: 0.7359307359307359

Confusion Matrix:

```
[[ 50  30]  
 [ 31 120]]
```

In [ ]: